

Unit 2: Uninformed, Informed & Local Search Strategies

Handwritten Student Notebook

Syllabus Topics Covered:

OFFICIAL SYLLABUS TOPICS COVERED Uninformed (Blind) Search: Breadth-First Search (BFS), Uniform-Cost Search (UCS), Depth-First Search (DFS), Depth-Limited Search (DLS), Iterative Deepening DFS (IDDFS), Bidirectional Search. Search Evaluation Criteria: Completeness, Time com...

Table of Contents

Uninformed (Blind) Search Strategies	p. 2
Comparative Evaluation Matrix of Search Strategies	p. 3
Informed Search & Heuristic Functions	p. 4
A*A* Search: Conditions for Optimality (Admissibility & Consistency)	p. 5
Heuristics for the 8-Puzzle & Heuristic Dominance	p. 9
Local Search Algorithms & Optimization	p. 10
Simulated Annealing & Genetic Algorithms	p. 12
One-Page Rapid Revision Sheet	p. 14
Important Examination Questions	p. 15
Authentic Solved Questions from Past Ramaiah CIE & SEE Papers	p. 18
Verified High-Yield Practice Problems	p. 25

Uninformed (Blind) Search Strategies

§ 1

Uninformed search (also termed *blind search*) refers to strategies that have no additional information about states beyond that provided in the problem definition. They can generate successors and distinguish a goal state from a non-goal state, but cannot estimate how close a non-goal state is to the destination.

1. Breadth-First Search (BFS)

BFS expands the shallowest unexpanded node first using a **First-In First-Out (FIFO) queue** frontier. All nodes at depth d are expanded before any nodes at depth $d + 1$.

- **Completeness:** Yes (if branching factor b is finite).
- **Time Complexity:** $O(b^d)$, where b is the branching factor and d is the depth of the shallowest goal.
- **Space Complexity:** $O(b^d)$ — every generated node must remain in memory.
Memory is the fatal bottleneck of BFS!
- **Optimality:** Yes, if all step costs are equal; otherwise not necessarily optimal.

2. Uniform-Cost Search (UCS / Dijkstra on Trees)

UCS expands the node n with the lowest path cost $g(n)$ from the start node using a **Priority Queue** ordered by $g(n)$.

- **Goal Test Timing:** The goal test is applied when a node is *selected for expansion* (popped from the priority queue), **NOT** when it is first generated!
- **Completeness:** Yes, provided every step cost is strictly positive ($c(s, a, s') \geq \epsilon > 0$).
- **Time & Space Complexity:** $O(b^{1+\lceil C^*/\epsilon \rceil})$, where C^* is the cost of the optimal solution and ϵ is the minimal step cost.
- **Optimality:** Yes! UCS is guaranteed to find the cheapest path for arbitrary non-negative step costs.

3. Depth-First Search (DFS)

DFS expands the deepest unexpanded node first using a **Last-In First-Out (LIFO) stack** frontier.

- **Completeness:** No in infinite-depth spaces or spaces with loops (unless cycle-checking graph search is used); Yes in finite state spaces.
- **Time Complexity:** $O(b^m)$, where m is the maximum depth of the state space ($m \gg d$).
- **Space Complexity:** $O(bm)$ — only needs to store a single path from the root to a leaf plus unexpanded sibling nodes. **Linear memory advantage!**
- **Optimality:** No — it can return a deeply nested suboptimal goal before exploring a shallow optimal one.

4. Iterative Deepening Depth-First Search (IDDFS)

IDDFS solves the dilemma between BFS (optimal, but exponential memory) and DFS (linear memory, but non-optimal and incomplete). It repeatedly runs **Depth-Limited Search (DLS)** with increasing depth limits ($l = 0, 1, 2, 3, \dots, d$).

WHY IDDFS IS THE PREFERRED UNINFORMED SEARCH METHOD

Although it appears wasteful to re-generate nodes from earlier depths, the bottom layer dominates the expansion cost in exponential trees. The number of generated nodes in IDDFS is:

$$N(\text{IDDFS}) = (d)b^1 + (d-1)b^2 + \dots + (1)b^d = O(b^d)$$

For $b = 10, d = 5$, BFS generates 111,111 nodes while IDDFS generates 123,450 nodes (only ~11% overhead!). IDDFS combines the **completeness and optimality of BFS** with the **modest $O(bd)$ linear memory storage of DFS**.

QUICK RECALL

- BFS: FIFO queue, $O(b^d)$ time & space, optimal for unit costs. Space is the killer.
- UCS: Priority queue ordered by $g(n)$, goal test on POP, optimal for non-negative costs.
- DFS: LIFO stack, $O(bm)$ linear space, not optimal, incomplete on infinite trees.
- IDDFS: Runs DLS with limit $0, 1, 2, \dots, d$. Combines BFS optimality with DFS $O(bd)$ linear space!

Comparative Evaluation Matrix of Search Strategies

§ 2

The standard evaluation metrics in AI are: **Completeness** (does it guarantee finding a solution?), **Time Complexity** (number of nodes generated), **Space Complexity** (maximum nodes stored in memory), and **Optimality** (does it find the lowest-cost solution?).

Criterion	Breadth-First (BFS)	Uniform-Cost (UCS)	Depth-First (DFS)	Depth-Limited (DLS)	Iterative Deepening (IDDFS)
Frontier Data Structure	FIFO Queue	Priority Queue on $g(n)$	LIFO Stack	LIFO Stack with limit l	LIFO Stack iterated $l = 0 \dots d$
Complete?	Yes (if b finite)	Yes (if step cost $\geq \epsilon > 0$)	No (infinite spaces / cycles)	No (if $l < d$)	Yes (if b finite)
Time Complexity	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space Complexity	$O(b^d)$ [Exponential!]	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$ [Linear!]	$O(bl)$ [Linear!]	$O(bd)$ [Linear!]
Optimal?	Yes (if unit costs)	Yes	No	No	Yes (if unit costs)

Notation: b = branching factor; d = depth of shallowest optimal goal; m = maximum depth of search tree; l = depth limit cutoff; C^* = optimal solution cost; ϵ = minimum step cost.

QUICK RECALL

- BFS space is $O(b^d)$ (exponential). DFS space is $O(bm)$ (linear).
- IDDFS space is $O(bd)$ (linear) and time is $O(b^d)$.
- UCS is optimal for general non-negative step costs.

Informed Search & Heuristic Functions

§ 3

Informed search strategies utilize problem-specific knowledge beyond the problem definition via a **heuristic function** $h(n)$, enabling the search to prune barren paths and focus exploration directly toward the goal.

HEURISTIC FUNCTION DEFINITION

A **heuristic function** $h(n)$ estimates the cost of the cheapest path from the state at node n to a goal state:

$$h(n) = \text{Estimated cost of shortest path from } n \text{ to Goal}$$

Property: If n is a goal node itself, then $h(n) = 0$.

Greedy Best-First Search vs A* Search

Search Algorithm	Evaluation Function $f(n)$	Behavioral Logic	Completeness & Optimality
Greedy Best-First Search	$f(n) = h(n)$	Expands the node that appears closest to the goal. Short-sighted: completely ignores past path cost $g(n)$.	Incomplete (can get trapped in loops); Non-optimal . Worst-case time/space $O(b^m)$.
A* Search	$f(n) = g(n) + h(n)$	Combines past cost $g(n)$ (cost to reach n) with estimated future cost $h(n)$. $f(n)$ is estimated cost of cheapest solution passing through n .	Complete and Optimal under appropriate heuristic conditions (Admissibility / Consistency).

QUICK RECALL

- $h(n)$: estimated cost from n to goal ($h(\text{goal}) = 0$).
- Greedy Best-First: $f(n) = h(n)$ (fast but non-optimal and incomplete).
- A* Search: $f(n) = g(n) + h(n)$ (optimal and complete).

A* Search: Conditions for Optimality (Admissibility & Consistency)**§ 4**

A* is the most widely celebrated search algorithm in computer science. However, its guarantee of optimality relies strictly on mathematical conditions governing the heuristic function.

1. ADMISSIBILITY (REQUIRED FOR TREE SEARCH)

An admissible heuristic is one that **never overestimates** the true cost to reach the goal. That is, $h(n)$ is optimistic:

$$h(n) \leq h^*(n) \quad \forall n$$

where $h^*(n)$ is the true, optimal cost to reach the nearest goal from n .

Theorem: If $h(n)$ is admissible, A^* using Tree-Search is **guaranteed to return an optimal solution**.

2. CONSISTENCY / MONOTONICITY (REQUIRED FOR GRAPH SEARCH)

A heuristic $h(n)$ is **consistent** (or monotonic) if, for every node n and every successor n' generated by action a :

$$h(n) \leq c(n, a, n') + h(n')$$

Geometric Intuition: This is a form of the **Triangle Inequality**. The estimated cost from n to the goal cannot exceed the step cost to step n' plus the estimated cost from n' to the goal.

Theorem: If $h(n)$ is consistent, A^* using Graph-Search is **guaranteed to be optimal**. Furthermore, the first time A^* expands any node n , the path to n is guaranteed to be optimal; hence nodes placed on the Closed list *never need to be reopened!*

RELATIONSHIP BETWEEN ADMISSIBILITY & CONSISTENCY

Every consistent heuristic is admissible ($Consistency \implies Admissibility$). While an admissible heuristic can occasionally be inconsistent in rare contrived graphs, virtually all natural heuristics derived from relaxed problems are automatically consistent.

WORKED TRACE · FULLY VERIFIED PYTHON SCRIPT [10 MARKS · EXAM CORE]**Step-by-Step A^* Graph Search Trace**

Trace A^* Graph Search from Start **S** to Goal **G**. The nodes have consistent heuristics h and directed edge costs c :

- Heuristics: $h(S) = 7, h(A) = 6, h(B) = 4, h(C) = 2, h(D) = 1, h(G) = 0$
- Edge costs: $(S, A) = 1, (S, B) = 4, (A, B) = 2, (A, C) = 5, (B, C) = 2, (C, D) = 2, (C, G) = 4, (D, G) = 1$

Step-by-Step Expansion Trace Table:

Iteration	Popped Node u	$g(u)$	$h(u)$	$f(u)$	Priority Queue (Open List) after Successor Expansion	Closed Set
Init	—	—	—	—	$[S(0 + 7 = 7)]$	$\emptyset$
1	S	0	7	7	Expand $S \rightarrow A(1 + 6 = 7), B(4 + 4 = 8)$. Queue: $[A(7), B(8)]$	$\{S\}$
2	A	1	6	7	Expand $A \rightarrow B(3 + 4 = 7), C(6 + 2 = 8)$. Queue: $[B(7), B(8), C(8)]$	$\{S, A\}$
3	B (path $S \rightarrow A \rightarrow B$)	3	4	7	Expand $B \rightarrow C(5 + 2 = 7)$. Queue: $[C(7), B(8), C(8)]$	$\{S, A, B\}$
4	C (path $S \rightarrow A \rightarrow B \rightarrow C$)	5	2	7	Expand $C \rightarrow D(7 + 1 = 8), G(9 + 0 = 9)$. Queue: $[B(8), C(8), D(8), G(9)]$	$\{S, A, B, C\}$
5	B (path $S \rightarrow B$)	4	4	8	B is already in Closed set! Discarded.	$\{S, A, B, C\}$
6	C (path $S \rightarrow A \rightarrow C$)	6	2	8	C is already in Closed set! Discarded.	$\{S, A, B, C\}$
7	D (path $S \rightarrow A \rightarrow B \rightarrow C \rightarrow D$)	7	1	8	Expand $D \rightarrow G(8 + 0 = 8)$. Queue: $[G(8), G(9)]$	$\{S, A, B, C, D\}$
8	G (Goal Popped!)	8	0	8	Goal reached on pop! Search halts.	$\{S, A, B, C, D, G\}$

Optimal Solution Found by A^* :

$S \rightarrow A \rightarrow B \rightarrow C \rightarrow D \rightarrow G$ with Optimal Cost $C^* = 8$

Notice that the alternative path $S \rightarrow A \rightarrow B \rightarrow C \rightarrow G$ has cost $5 + 4 = 9$. A^* correctly discovered the superior detour through D with cost 8!

QUICK RECALL

- Admissibility ($h(n) \leq h^*(n)$): never overestimates. Guarantees optimality in tree search.
- Consistency ($h(n) \leq c(n, a, n') + h(n')$): triangle inequality. Guarantees optimality in graph search without reopening closed nodes.
- Goal test in A^* must be applied when a node is **popped**, not generated.

 **Interactive Tool Available Online:** ⚡ Interactive Search Strategy & Simulated Annealing Simulator. Try the interactive step-by-step calculator and visualizer on the live site notes page.

Heuristics for the 8-Puzzle & Heuristic Dominance

§ 5

A standard method for generating admissible heuristics is to compute the exact cost of a solution to a **relaxed problem** (a problem with fewer restrictions on the actions).

Two Classical 8-Puzzle Heuristics

1. $h_1(n)$ — Number of Misplaced Tiles:

Count the number of tiles that are not in their goal position. (If a tile can move anywhere instantly, each misplaced tile needs at least 1 move). Clearly admissible: $h_1(n) \leq h^*(n)$.

2. $h_2(n)$ — Manhattan Distance (City Block Distance):

The sum of the horizontal and vertical distances of each tile from its target goal position:

$$h_2(n) = \sum_{i=1}^8 (|x_i - x_{i,\text{goal}}| + |y_i - y_{i,\text{goal}}|)$$

Admissible because each move can reduce the Manhattan distance of at most one tile by at most 1.

HEURISTIC DOMINANCE

If $h_2(n) \geq h_1(n)$ for all nodes n , and both are admissible, we say that h_2 **dominates** h_1 .

Consequence: A^* using h_2 will never expand more nodes than A^* using h_1 (except possibly for some nodes on the boundary where $f(n) = C^*$). A dominating heuristic translates directly to greater pruning power and faster search!

For the 8-puzzle: since Manhattan distance counts the minimum moves for each tile to reach home and each misplaced tile is at least 1 step away, $h_2(n) \geq h_1(n)$ always holds. **Manhattan distance strictly dominates misplaced tiles.**

QUICK RECALL


- Relaxed problem: removing constraints yields admissible heuristics.
- 8-puzzle: h_1 = misplaced tiles; h_2 = Manhattan distance ($\sum |\Delta x| + |\Delta y|$).
- h_2 dominates $h_1 \iff h_2(n) \geq h_1(n) \implies h_2$ expands fewer nodes.

Local Search Algorithms & Optimization

§ 6

In many problems (such as the 8-queens, VLSI circuit layout, or factory scheduling), the path to the goal is completely irrelevant; only the **final goal state itself matters**. **Local search algorithms** operate using a single current state (rather than multiple paths) and generally move only to neighbors of that state, requiring $O(1)$ constant memory.

State Space Landscape & Local Search Hazards



Hill-Climbing Search (Greedy Local Search)

Hill-climbing simply loops continuously in the direction of increasing value (uphill). It terminates when it reaches a peak where no neighbor has a higher value.

THREE CLASSIC FAILURE MODES OF HILL CLIMBING

1. **Local Maxima:** A peak that is higher than each of its neighbors, but lower than the global maximum. The algorithm terminates prematurely.
2. **Ridges:** A sequence of local maxima where the slope gently rises along a diagonal ridge, but the allowable orthogonal moves all lead downhill.
3. **Plateaux:** A flat area of the state space landscape where all neighbors have the identical value. The algorithm wanders aimlessly in a random walk.

Remedy — Random-Restart Hill Climbing: Conducts a series of hill-climbing searches from randomly generated initial states until a goal is found. *Theorem:* If each hill-climb has probability p of success, expected number of restarts is $1/p$. Trivially complete with probability 1!

QUICK RECALL

- Local search operates on complete-state representations with $O(1)$ memory.
- Hill climbing moves greedily uphill; gets stuck on local maxima, ridges, and plateaux.
- Random-Restart Hill Climbing guarantees finding global optimum given sufficient restarts.

Simulated Annealing & Genetic Algorithms

§ 7

To escape local maxima, an algorithm must occasionally make **downhill moves**. Two primary nature-inspired stochastic optimization methods achieve this: **Simulated Annealing** (thermodynamics) and **Genetic Algorithms** (biological evolution).

1. Simulated Annealing

Inspired by the metallurgical process of annealing (heating metal to high temperatures and cooling it slowly to form a crystal with minimal energy defects).

SIMULATED ANNEALING ALGORITHM

At each step, choose a random neighbor state $next$. Let $\Delta E = Value(next) - Value(current)$:

- If $\Delta E > 0$ (the move improves objective value), **always accept the move**.
- If $\Delta E \leq 0$ (the move is worse / downhill), accept the move with probability:

$$P = e^{\frac{\Delta E}{T}}$$

where $T > 0$ is a control parameter called the **Temperature**, governed by a cooling schedule.

WORKED NUMERICAL · FULLY VERIFIED PASS**[5 MARKS · EXAM CORE]****Downhill Acceptance Probabilities in Simulated Annealing**

In simulated annealing, an agent encounters a bad downhill move with energy penalty $\Delta E = -5$. Compute the acceptance probability $P = e^{\Delta E/T}$ at temperatures $T = 1000, 100, 10, 1$.

Temperature T	Exponent $\Delta E/T$	Probability $P = e^{\Delta E/T}$	Physical / Algorithmic Behavior
$T = 1000$ K (Hot)	$-5/1000 = -0.005$	$e^{-0.005} \approx$ 0.9950 (99.50%)	Almost all downhill moves accepted; behaves like random walk to escape traps.
$T = 100$ K (Warm)	$-5/100 = -0.05$	$e^{-0.05} \approx$ 0.9512 (95.12%)	High exploration, readily navigates ridges.
$T = 10$ K (Cooling)	$-5/10 = -0.5$	$e^{-0.5} \approx$ 0.6065 (60.65%)	Moderate exploration; bad moves increasingly penalized.
$T = 1$ K (Cold)	$-5/1 = -5.0$	$e^{-5.0} \approx$ 0.0067 (0.67%)	Freezes into pure greedy hill climbing; downhill moves rejected.

Conclusion: If the cooling schedule lowers T slowly enough, the algorithm is proven to converge to the global optimum with probability 1!

2. Genetic Algorithms (GAs)

A variant of stochastic beam search where successor states are generated by combining two parent states (crossover) rather than modifying a single state.

1

Representation: A state is encoded as a string over a finite alphabet (a chromosome / string of digits, e.g. 8-queens encoded as an 8-digit string indicating the row position of queen in each column).

2

Fitness Function: Evaluates the quality of each state (e.g. number of non-attacking queen pairs, maximum = 28).

3

Selection: Pairs of individuals are chosen for reproduction with probability proportional to their fitness (Roulette-wheel selection).

4

Crossover: A random crossover point is chosen. Offspring inherit the first segment from parent 1 and the remainder from parent 2.

5

Mutation: Each bit/digit is randomly mutated with a small independent probability (e.g. 1%) to maintain genetic diversity.

QUICK RECALL

- Simulated Annealing accepts downhill moves with probability $P = e^{\Delta E/T}$. High T = random walk; Low T = greedy hill climb.
- Genetic Algorithms: Population, Fitness proportional selection, Crossover, Mutation.
- 8-queens can be encoded as an 8-character string with fitness = non-attacking pairs.

One-Page Rapid Revision Sheet

★ REVISION

Blind Search Complexities

- BFS: FIFO, $O(b^d)$ time, $O(b^d)$ space, optimal for unit cost.
- UCS: Priority queue $g(n)$, goal test on pop, optimal for general costs.
- DFS: LIFO, $O(b^m)$ time, $O(bm)$ linear space, non-optimal.
- IDDFS: $O(b^d)$ time, $O(bd)$ linear space, optimal and complete!

A* Search Checklist

- $f(n) = g(n) + h(n)$.
- Admissibility: $h(n) \leq h^*(n)$ (optimistic). Optimal for tree search.
- Consistency: $h(n) \leq c(n, a, n') + h(n')$ (triangle inequality). Optimal for graph search.
- *Consistency* $\implies$ *Admissibility*.

8-Puzzle Heuristics

- h_1 : misplaced tiles ($h_1 \leq h^*$).
- h_2 : Manhattan distance ($\sum |\Delta x| + |\Delta y|$).
- h_2 dominates h_1 ($h_2(n) \geq h_1(n)$) $\implies$ expands fewer nodes.

Local Search & Annealing

- Hill climbing: gets stuck on local maxima, ridges, plateaux.
- Simulated Annealing: accepts downhill moves with $P = e^{\Delta E/T}$.
- High $T \implies P \approx 1$; Low $T \implies P \approx 0$.
- Genetic Algorithms: Selection, Crossover, Mutation.

Important Examination Questions

? QUESTIONS

2-Mark Questions (Short Answer)

Practice Problem:

2M Define an admissible heuristic function.

Ans

Answer: A heuristic function $h(n)$ is admissible if it never overestimates the true cost to reach the goal from node n , i.e., $h(n) \leq h^*(n)$ for all nodes n , where $h^*(n)$ is the true optimal cost to the goal.

Practice Problem:

2M Why is IDDFS preferred over BFS and DFS?

Ans

Answer: IDDFS combines the best benefits of both algorithms: it achieves the completeness and optimality of BFS while requiring only the modest linear memory storage $O(bd)$ of DFS.

Practice Problem:

2M What is the acceptance probability formula in simulated annealing?

Ans

Answer: $P = e^{\Delta E/T}$, where $\Delta E < 0$ is the change in objective energy/value for a downhill move, and $T > 0$ is the temperature parameter.

5-Mark Questions (Medium Length)

Practice Problem:

5M Prove that consistency of a heuristic implies admissibility.

Ans

Answer:

1. By consistency, $h(n) \leq c(n, a, n') + h(n')$.
2. Applying this along the sequence of nodes on the optimal path from n to goal G :
$$h(n) \leq c(n, a_1, n_1) + h(n_1) \leq c(n, a_1, n_1) + c(n_1, a_2, n_2) + \dots + h(G)$$
3. Since G is a goal node, $h(G) = 0$.
4. The sum of the step costs along the optimal path is exactly $h^*(n)$.
5. Therefore, $h(n) \leq h^*(n)$, proving $h(n)$ is admissible.

Practice Problem:

5M Explain the failure modes of hill-climbing search.

Ans

Answer:

- **Local Maxima:** Peaks that are higher than all neighbors but lower than the global maximum; search stops prematurely.
- **Ridges:** Sequences of local maxima where the slope gently inclines along a ridge, but individual orthogonal moves lead downhill, causing oscillation.
- **Plateaux:** Flat areas where all neighboring states have the same value; search degenerates into a random walk.

10-Mark Questions (Long Answer / Comprehensive)



Practice Problem:

10M Explain A^* search with an example. Prove its optimality for tree search and graph search.

Ans

Answer Summary: Detail evaluation function $f(n) = g(n) + h(n)$. Present the admissibility proof for tree search (showing that a suboptimal goal G_2 cannot be popped before an optimal goal G), and the consistency proof for graph search ($f(n)$ non-decreasing along any path). Provide the step-by-step trace from Section 4.

Authentic Solved Questions from Past Ramaiah CIE & SEE Papers

■ EXAM ARCHIVE

The following questions are transcribed directly from authentic internal assessment and semester-end examination papers in notes/ai/practice/ai-cie-1-and-2.pdf (Ramaiah Continuous Internal Evaluation - I, October 29, 2025, Course Code ISE552):

Ramaiah CIE-1 · October 29, 2025 · Question 1.b [4 Marks]

"Describe Bidirectional Search. Trace BFS and DFS traversal for the given graph from source S to goal G ."

 Practice Problem:

Show Complete Step-by-Step Model Answer

Ans

1. Principle of Bidirectional Search:

Bidirectional search replaces a single forward search from initial state S to goal G with **two simultaneous searches**: one forward from S , and one backward from G . Search terminates successfully when the two search frontiers intersect in the middle.

- If the branching factor is b and optimal goal distance is d , standard BFS expands $O(b^d)$ nodes.
- Bidirectional BFS runs two trees of depth $d/2$, expanding $O(b^{d/2} + b^{d/2}) = O(2b^{d/2})$ nodes.
- **Dramatic Speedup:** For $b = 10$ and $d = 6$, standard BFS expands $10^6 = 1,000,000$ nodes, whereas bidirectional BFS expands only $2 \times 10^3 = 2,000$ nodes — a 500times reduction!

2. Comparative Traversal Trace on Benchmark Graph:

Search Strategy	Data Structure	Order of Node Expansion	Solution Path Returned	Optimality Verdict
BFS	FIFO Queue	Level 0: $S \rightarrow$ Level 1: $A, B \rightarrow$ Level 2: $C, D \rightarrow$ Level 3: G	$SoAoCoG$ (Cost = $1 + 5 + 4 = 10$)	Optimal in number of hops (3 hops), but suboptimal in path cost!
DFS	LIFO Stack	Deepest path first: $SoAoDoG$	$SoAoDoG$ (Cost = $1 + 12 + 1 = 14$)	Not optimal (falls into deep branches; high cost 14).
		Forward: $So\{A, B\};$		Guaranteed optimal when

Ramaiah CIE-1 · October 29, 2025 · Question 2.b [6 Marks]

"Demonstrate A* search on the 8-puzzle problem using misplaced tiles heuristic $h_1(n)$. Show the search tree expansion step-by-step with $f(n) = g(n) + h(n)$ values."

 Practice Problem:

Show Complete Step-by-Step Model Answer

Ans

1. Problem Formulation:

- Evaluation function: $f(n) = g(n) + h_1(n)$, where $g(n)$ is the depth (number of moves from start) and $h_1(n)$ is the number of misplaced tiles (excluding the blank).
- Goal state: Tiles 1 to 8 arranged in row-major order: $\text{ext}\{\text{Goal}\} = \begin{bmatrix} 1 & 2 & 3 & \text{blank} \\ 8 & 4 & 7 & 6 \\ 5 & \end{bmatrix}$
- Start state: $\text{ext}\{\text{Start}\} = \begin{bmatrix} 2 & 8 & 3 & 1 \\ \text{blank} & 4 & 7 & 6 \\ 5 & \end{bmatrix}$

2. Step-by-Step Search Tree Expansion:

1. Depth 0 (Root Node):

- Comparing Start to Goal: Tiles 1, 2, and 8 are misplaced. Tiles 3, 4, 5, 6, 7 are in goal positions.
- $g(\text{extRoot}) = 0$, $h_1(\text{extRoot}) = 3$ (or 4 depending on numbering). Here $h_1 = 4$.
- $f(\text{extRoot}) = 0 + 4 = 4$.

2. Depth 1 Successors (Blank Moves):

- **Action UP (slide 8 down):** Tiles misplaced = 5. $f = 1 + 5 = 6$.
- **Action LEFT (slide 1 right):** Moving tile 1 places it closer to goal. Tiles misplaced = 3. $f = 1 + 3 = 4$ ★ (Minimum f).
- **Action RIGHT (slide 4 left):** Tiles misplaced = 5. $f = 1 + 5 = 6$.
- **Action DOWN (slide 6 up):** Tiles misplaced = 5. $f = 1 + 5 = 6$.

3. OPEN Priority Queue Selection:

- Node LEFT has the lowest $f(n) = 4$, so it is chosen for immediate expansion.
- All three sibling nodes with $f(n) = 6$ are preserved in the OPEN list without expansion.

4. Depth 2 Successor from Node LEFT:

- From $\begin{bmatrix} 2 & 8 & 3 & \text{blank} \\ 1 & 4 & 7 & 6 \\ 5 & \end{bmatrix}$, moving UP slides 2 into its goal position, reducing

Ramaiah CIE-1 · October 29, 2025 · Question 3.b [4 Marks]

"Compare Uninformed Search with Uniform Cost Search (UCS). Explain their time and space complexities."

 Practice Problem:

Show Complete Step-by-Step Model Answer

Ans

1. Distinction:

- **Standard Uninformed Search (BFS):** Expands nodes in FIFO order based purely on *tree depth* (number of hops from root). Assumes all action step costs are identical ($c = 1$).
- **Uniform Cost Search (UCS / Dijkstra's Search):** Expands nodes in order of *accumulated path cost* $g(n)$ using a priority queue. Works on arbitrary, non-uniform positive edge step costs $c(s, a, s') \geq \epsilon > 0$.

2. Complexity Derivation:

- Let C^* be the optimal path cost to the goal, and let ϵ be the minimum step cost between any two connected nodes.
- In the worst case, the search may explore nodes along paths up to depth $\lfloor C^* / \epsilon \rfloor$.
- **Time Complexity:** $\mathbf{O(b^{1 + \lfloor C^* / \epsilon \rfloor})}$ If all step costs are equal ($\epsilon = c$), then $C^* / \epsilon = d$, and time complexity simplifies to $O(b^d)$, matching standard BFS.
- **Space Complexity:** $\mathbf{O(b^{1 + \lfloor C^* / \epsilon \rfloor})}$ UCS must store all explored and frontier nodes in the priority queue to maintain global minimum-cost ordering.

3. Optimality Condition:

UCS is guaranteed optimal **if and only if all step costs are strictly positive** ($c \geq \epsilon > 0$). If zero-cost or negative-cost cycles exist, UCS can enter infinite loops without advancing towards the goal.

Ramaiah SEE Archive · Model Question Paper [8 Marks]

"Prove the optimality of A^* search for: (a) Tree search under an admissible heuristic, (b) Graph search under a consistent (monotonic) heuristic."

 Practice Problem:

Show Complete Step-by-Step Model Answer

Ans

(a) Optimality Proof for Tree Search (A^* with Admissible Heuristic):

1. Suppose G_2 is a suboptimal goal node on the OPEN list, so $g(G_2) > C^*$, where C^* is the optimal solution cost. Since G_2 is a goal node, $h(G_2) = 0$, giving $f(G_2) = g(G_2) > C^*$.
2. Let n be an unexpanded node on the OPEN list that lies on an optimal path to the optimal goal G .
3. Because $h(n)$ is admissible ($h(n) \leq h^*(n)$):

$$f(n) = g(n) + h(n) \leq g(n) + h^*(n) = C^*$$

4. Combining both inequalities:

$$f(n) \leq C^* < f(G_2)$$

5. Since $f(n) < f(G_2)$, the priority queue will pop node n before G_2 . By induction, all nodes on the optimal path are expanded before G_2 can ever be popped. Therefore, A^* tree search never selects a suboptimal goal.

□ lacksquare

(b) Optimality Proof for Graph Search (A^* with Consistent Heuristic):

1. A heuristic is **consistent (monotonic)** if for every node n and every successor n' generated by action a :

$$h(n) \leq c(n, a, n') + h(n')$$

2. Adding $g(n)$ to both sides:

$$g(n) + h(n) \leq g(n) + c(n, a, n') + h(n')$$

$$f(n) \leq g(n') + h(n') = f(n')$$

3. Thus, $f(n)$ is **monotonically non-decreasing along any path**.
4. Consequently, whenever A^* selects a node n for expansion from the OPEN list, the path found to n is guaranteed to be optimal. No previously expanded node on the CLOSED list ever needs to be reopened,

Ramaiah AI Question Bank · Stochastic Optimization [6 Marks]

"Explain the operation of Simulated Annealing. Detail how the cooling schedule and Boltzmann probability $P = e^{\Delta E/T}$ prevent the search from getting trapped in local maxima."

**Practice Problem:**

Show Complete Step-by-Step Model Answer

Ans**1. Inspiration & Algorithmic Framework:**

Adapted from metallurgical annealing (gradual cooling of molten metals to reach low-energy crystalline states), Simulated Annealing escapes local maxima by combining greedy hill-climbing with stochastic exploration.

2. Downhill Move Acceptance Rule:

- At state s , pick a random neighboring state s' . Compute energy difference $\Delta E = E(s') - E(s)$.
- If $\Delta E > 0$ (uphill move towards improvement), the move is **always accepted**.
- If $\Delta E \leq 0$ (downhill move), the move is accepted probabilistically with **Boltzmann probability**:

$$P = e^{\Delta E/T}$$

where $T > 0$ is the current temperature parameter.

3. Role of Temperature (T) and Cooling Schedule:

- **Initial High Temperature ($T \gg 0$):** When T is very large, $\Delta E / T \approx 0 \implies P \approx e^0 = 1$. The search accepts almost all downhill moves, behaving like a random walk that freely explores the entire landscape without getting trapped.
- **Cooling Progression ($T_{k+1} = \alpha T_k, 0.8 \leq \alpha \leq 0.99$):** As T decreases, the probability P of accepting worse moves drops exponentially. The agent becomes increasingly selective, accepting only shallow downhill moves.
- **Freezing Phase ($T \rightarrow 0$):** As $T \rightarrow 0$, $P \rightarrow 0$. Downhill moves are strictly rejected, and the algorithm smoothly degenerates into standard greedy hill-climbing, locking cleanly into the nearest global peak.

Verified High-Yield Practice Problems

▽ PRACTICE

Test your problem-solving skills with these exam-caliber problems. All numerical calculations and algorithmic steps are verified against `audit/verify/ai/verify_ai_u2.py`:

PRACTICE 1 · BFS EXPONENTIAL MEMORY ANALYSIS

In a search tree with branching factor $b = 10$ and depth $d = 6$, calculate the number of nodes stored in memory by BFS if each node requires 100 bytes. What happens at depth $d = 8$?

Practice Problem:

Show Step-by-Step Solution

Ans

1. Memory Consumption at Depth $d = 6$:

$\text{Total Nodes Stored} \approx b^d = 10^6 = 1,000,000 \text{ nodes}$
 $\text{Memory} = 1,000,000 \times 100 \text{ bytes} = 100,000,000 \text{ bytes}$
 $\approx \mathbf{100 \text{ MB}}$

2. Memory Consumption at Depth $d = 8$:

$\text{Total Nodes Stored} \approx 10^8 = 100,000,000 \text{ nodes}$
 $\text{Memory} = 10^8 \times 100 \text{ bytes} = 10,000,000,000 \text{ bytes}$
 $\approx \mathbf{10 \text{ GB}}$

Takeaway: BFS memory scales exponentially ($O(b^d)$). In contrast, IDDFS achieves the exact same completeness and optimality at depth 8 while requiring only linear memory $O(bd) = 10 \times 8 \times 100 \text{ bytes} = \mathbf{8 \text{ KB}}$!

PRACTICE 2 · A* GRAPH SEARCH NUMERICAL TRACE

Trace A* search on the graph with nodes S, A, B, C, D, G , edge step costs $c(S, A) = 1, c(S, B) = 4, c(A, B) = 2, c(A, C) = 5, c(B, C) = 2, c(C, D) = 2, c(D, G) = 1, c(C, G) = 4$, and heuristic values $h(S) = 7, h(A) = 6, h(B) = 4, h(C) = 2, h(D) = 1, h(G) = 0$. Find the optimal path and cost.

 Practice Problem:

Show Step-by-Step Solution

Ans

Step-by-step priority queue expansion log (verified via `audit/verify/ai/verify_ai_u2.py`):

1. **Step 1:** Pop S ($g = 0, h = 7, f = 7$). Successors: A ($g = 1, h = 6, f = 7$), B ($g = 4, h = 4, f = 8$).
2. **Step 2:** Pop A ($g = 1, h = 6, f = 7$). Successors: B via A ($g = 3, h = 4, f = 7$), C via A ($g = 6, h = 2, f = 8$).
3. **Step 3:** Pop B ($g = 3, h = 4, f = 7$). Successor: C via B ($g = 5, h = 2, f = 7$).
4. **Step 4:** Pop C ($g = 5, h = 2, f = 7$). Successors: D ($g = 7, h = 1, f = 8$), G ($g = 9, h = 0, f = 9$).
5. **Step 5:** Pop D ($g = 7, h = 1, f = 8$). Successor: G ($g = 8, h = 0, f = 8$).
6. **Step 6:** Pop G ($g = 8, h = 0, f = 8$). Goal reached!

Optimal Path: SoAoBoCoDoG with Optimal Cost = 8.0.

PRACTICE 3 · HEURISTIC CONSISTENCY PROOF

Prove mathematically by mathematical induction that every consistent (monotonic) heuristic is necessarily admissible.

 Practice Problem:

Show Step-by-Step Solution

Ans

Proof:

1. By definition of consistency, for any state n and successor n' via action a :

$$h(n) \leq c(n, a, n') + h(n')$$

2. Let $n = n_0on_1on_2o \dots on_k = G$ be the optimal path from node n to goal G .

3. Applying the consistency inequality along this optimal chain:

$$\begin{aligned} h(n_0) &\leq c(n_0, a_1, n_1) + h(n_1) \\ h(n_0) &\leq c(n_0, a_1, n_1) + c(n_1, a_2, n_2) + h(n_2) \\ h(n_0) &\leq \sum_{i=0}^{k-1} c(n_i, a_{i+1}, n_{i+1}) + h(G) \end{aligned}$$

4. Since G is a goal state, $h(G) = 0$.
5. The sum of optimal step costs is by definition the true optimal cost $h^*(n_0)$:

$$h(n_0) \leq h^*(n_0) + 0 = h^*(n_0)$$

6. Therefore, $h(n) \leq h^*(n)$ for all nodes n , proving h is admissible.

□lacksquare

PRACTICE 4 · 8-PUZZLE HEURISTIC DOMINANCE

Calculate the Manhattan distance heuristic h_2 for the 8-puzzle start state where tile 5 is at $(1, 3)$ and goal is $(3, 2)$. Prove why h_2 strictly dominates h_1 (misplaced tiles).

 Practice Problem:

Show Step-by-Step Solution

Ans

1. Manhattan Calculation for Tile 5:

$$d_M(5) = |x_1 - x_2| + |y_1 - y_2| = |1 - 3| + |3 - 2| = 2 + 1 = 3 \text{ steps}$$

2. Dominance Proof ($h_2(n) \geq h_1(n)$):

- For any misplaced tile t , it must move at least 1 square to reach its goal position $\implies d_M(t) \geq 1$.
- For any tile already in its goal position, $d_M(t) = 0$.
- Therefore:

$$h_2(n) = \sum_{t=1}^8 d_M(t) \geq \sum_{t \text{ misplaced}} 1 = h_1(n)$$

- Since $h_2(n) \geq h_1(n)$ for all states n , h_2 **dominates** h_1 . As a consequence, A^* using h_2 will expand a subset of the nodes expanded by A^* using h_1 , saving substantial computation.

PRACTICE 5 · BOLTZMANN PROBABILITY CALCULATION

In Simulated Annealing with $\Delta E = -5.0$, compute the exact acceptance probability $P = e^{\Delta E/T}$ at $T = 1000\text{extK}$, 100extK , 10extK , and 1extK .

 Practice Problem:

Show Step-by-Step Solution

Ans

Computed via `audit/verify/ai/verify_ai_u2.py`:

- At $T = 1000\text{extK}$: $P = e^{\{-5 / 1000\}} = e^{\{-0.005\}}$ $\approx \mathbf{0.995012 \text{ ext } (99.50\%)}$.
- At $T = 100\text{extK}$: $P = e^{\{-5 / 100\}} = e^{\{-0.05\}}$ $\approx \mathbf{0.951229 \text{ ext } (95.12\%)}$.
- At $T = 10\text{extK}$: $P = e^{\{-5 / 10\}} = e^{\{-0.5\}}$ $\approx \mathbf{0.606531 \text{ ext } (60.65\%)}$.
- At $T = 1\text{extK}$: $P = e^{\{-5 / 1\}} = e^{\{-5.0\}}$ $\approx \mathbf{0.006738 \text{ ext } (0.67\%)}$.

Conclusion: At high T , worse moves are accepted with 99.5% probability (allowing escapes from local maxima). Near $T = 1$, the acceptance rate drops below 1%, forcing convergence.

PRACTICE 6 · UCS GOAL TEST TIMING & ZERO COSTS

Why must Uniform Cost Search apply the goal test when a node is selected for EXPANSION (popped from priority queue), rather than when it is GENERATED (inserted into queue)?

 Practice Problem:

Show Step-by-Step Solution

Ans**Answer:**

- If the goal test were applied when a node is *generated*, UCS could stop immediately upon discovering the first path to the goal.
- However, that first path might have a high step cost (e.g., *SoG* with cost 10). A much cheaper path might exist (e.g., *SoAoBoG* with total cost $1 + 1 + 1 = 3$), but its intermediate nodes have not yet been expanded!
- By testing for the goal only when a node is **popped for expansion**, UCS guarantees that all nodes with path cost less than $g(G)$ have already been expanded. This ensures that when G is popped, no cheaper path to G can possibly exist, preserving mathematical optimality.